

NETRA

Weekly Governance Brief

Generated: 2026-06-14 17:06 UTC

What This Means (Plain English)

In simple terms: NETRA reviewed 63 actions that AI agents tried to perform. Of these, 37 were dangerous and were automatically stopped before they could cause harm. This saved the organization an estimated Rs.14,755,067. 14 actions needed a human to make the final call. The system is working as intended - catching risky AI behavior before it happens.

Executive Summary (Technical)

NETRA processed 63 AI agent actions this period. 37 actions were blocked (58.7% block rate), preventing an estimated Rs.14,755,067 in potential damage. 14 actions were escalated for human review. The governance system delivered a 530.09x return on investment.

Key Metrics

Metric	Value	Metric	Value
Total Actions	63	Blocked	37
Approved	12	Escalated	14
Avg Risk Score	55.7	Max Risk Score	86.1
Damage Prevented	Rs.14,755,067	ROI Multiplier	530.09x

Top 5 Highest Risk Actions

Agent	Action	Amount	Risk	Customer	Reasoning
Inventory Bot	refund	Rs.75,000	86	cust_456	Policy violations detected: refund_limit
Inventory Bot	refund	Rs.75,000	86	cust_456	Policy violations detected: refund_limit
Inventory Bot	refund	Rs.75,000	86	cust_456	Policy violations detected: refund_limit
support_bot_v2.	refund	Rs.75,000	86	cust_456	Policy violations detected: refund_limit
Finance Bot	refund	Rs.75,000	85	cust_456	Policy violations detected: refund_limit

Recommendations

- HIGH AVG RISK: Review agent permissions and tighten policy thresholds.
- HIGH BLOCK RATE: Consider retraining agents or refining tool access.
- HIGH ESCALATION: Add more deterministic rules to reduce manual review load.
- Schedule quarterly governance policy review with stakeholders.
- Expand agent coverage to include Customer Service and HR bots.

